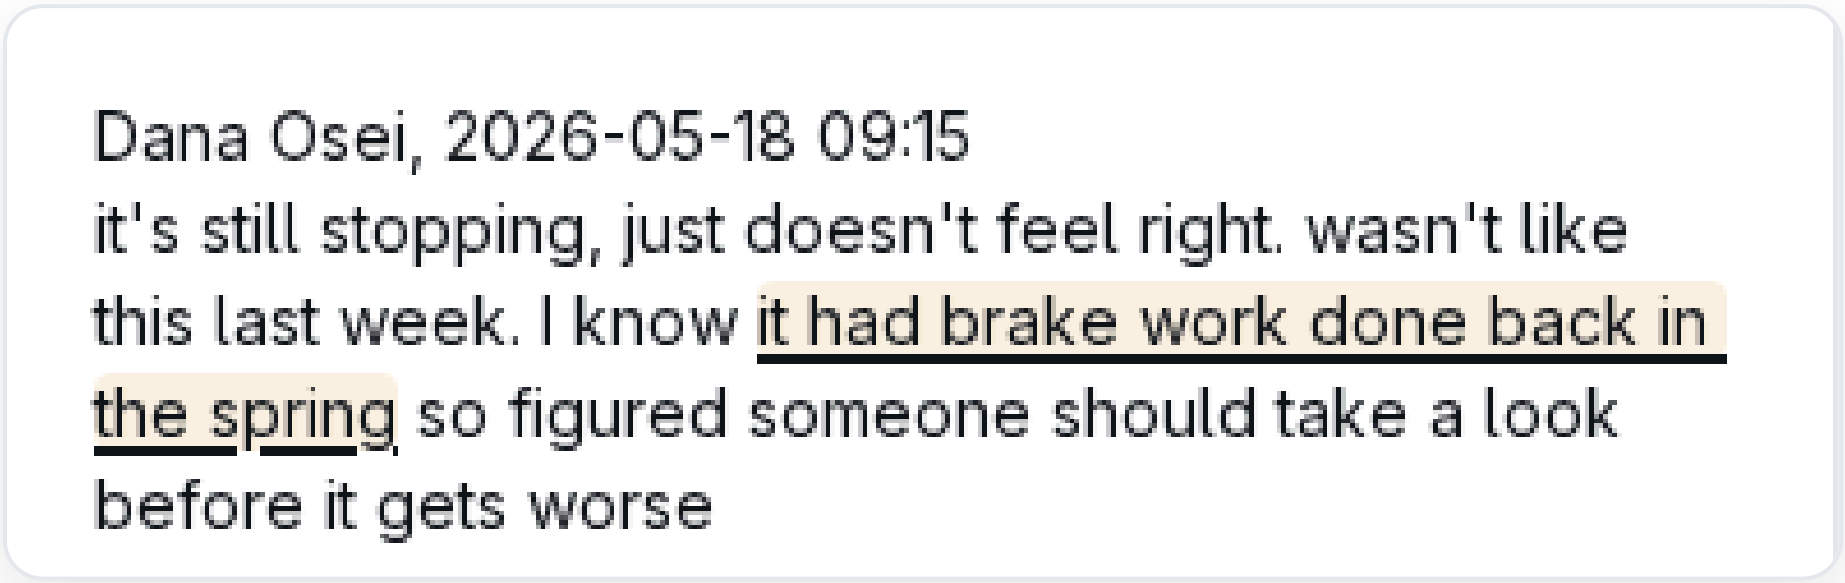
CAROUSEL

What a review queue should show a human

An extracted answer is not enough. These are the pieces one review item carries so a person can accept it, correct it or send it back.

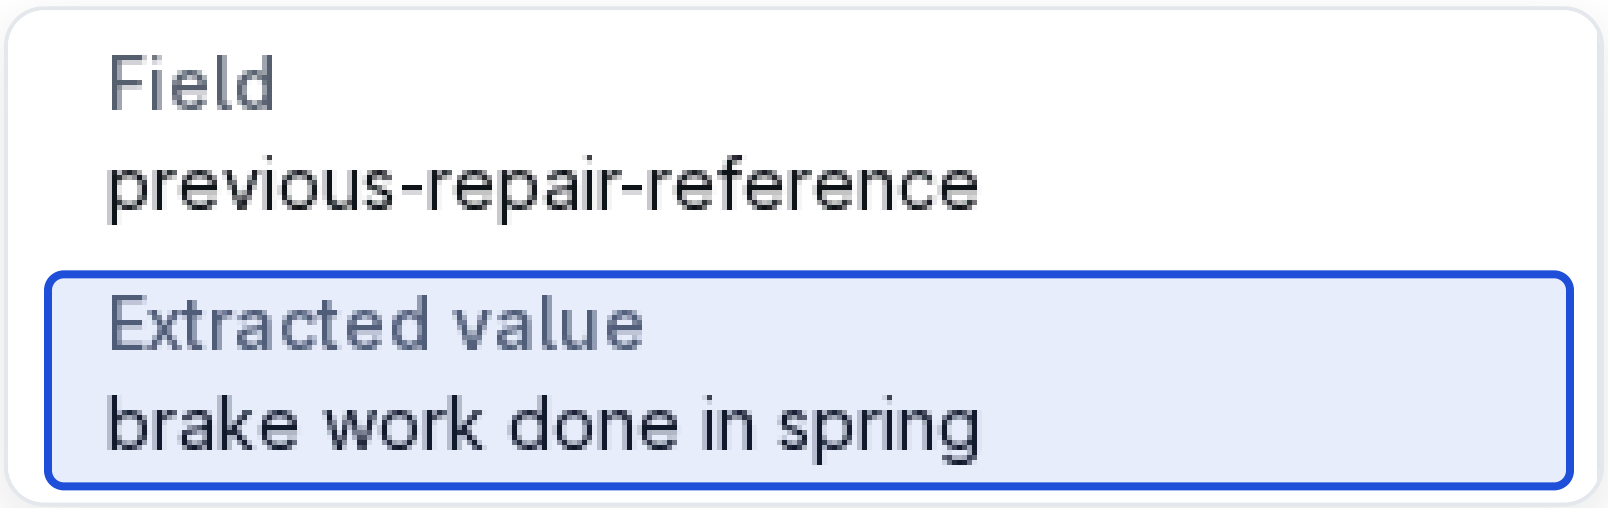
Show the source, not a summary

The reviewer sees the original chat message with the exact words behind the value underlined. They read the sentence itself, not a paraphrase.



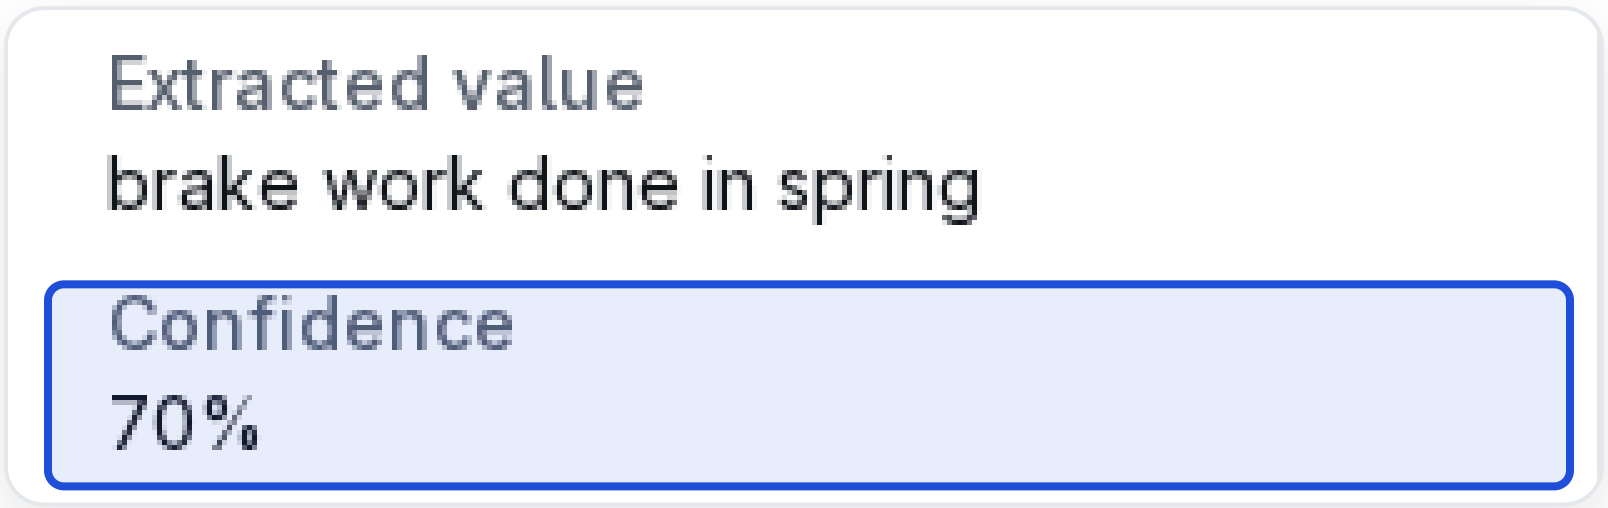
Show the value as extracted

Field: previous-repair-reference. Value: "brake work done in spring". The reviewer can judge whether the value really follows from the highlighted words.



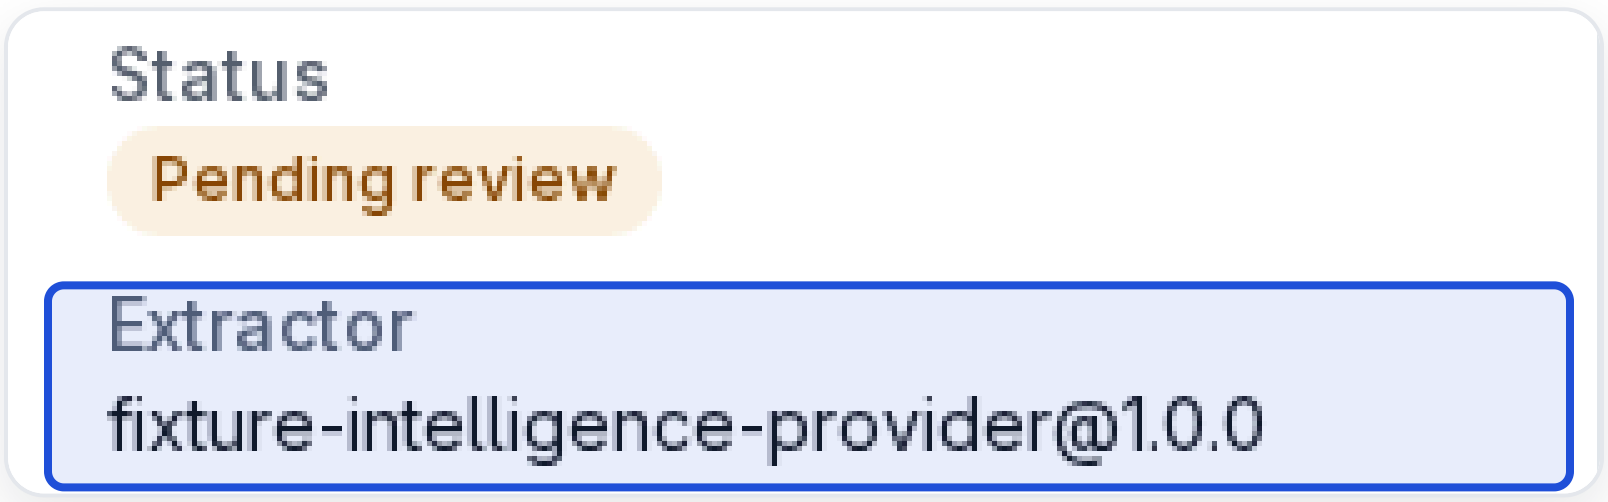
Show confidence as a number

This item shows 70%. In this demo that figure is authored fixture data, not a measured model score. A real model's confidence would need calibrating first.



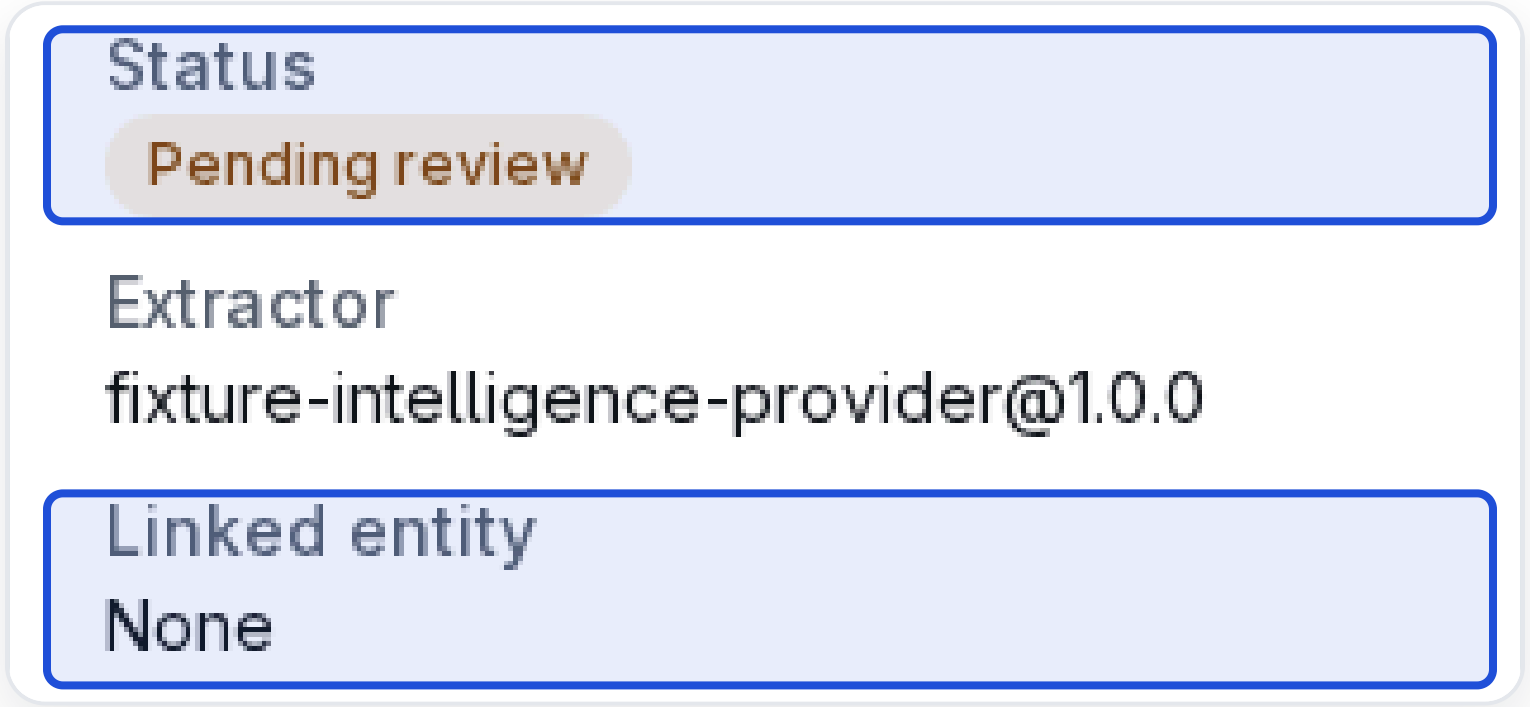
Show which extractor produced it

Extractor identity and version are on the item: fixture-intelligence-provider@1.0.0. If a result is wrong later, you know what produced it.



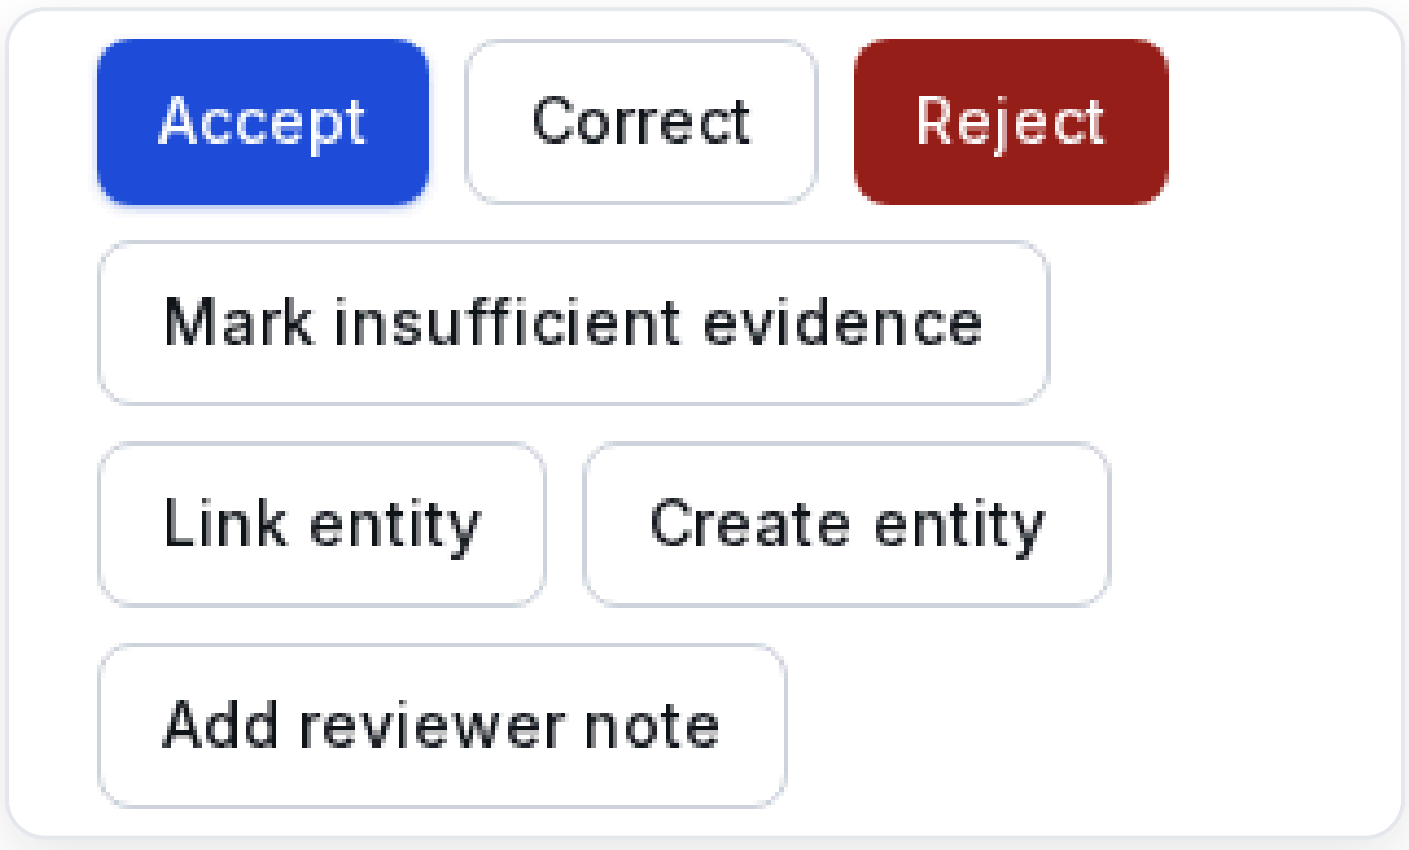
Show status and what it is linked to

Status: pending review. Linked entity: none. The reviewer can see the value is not yet attached to any asset, case or event.



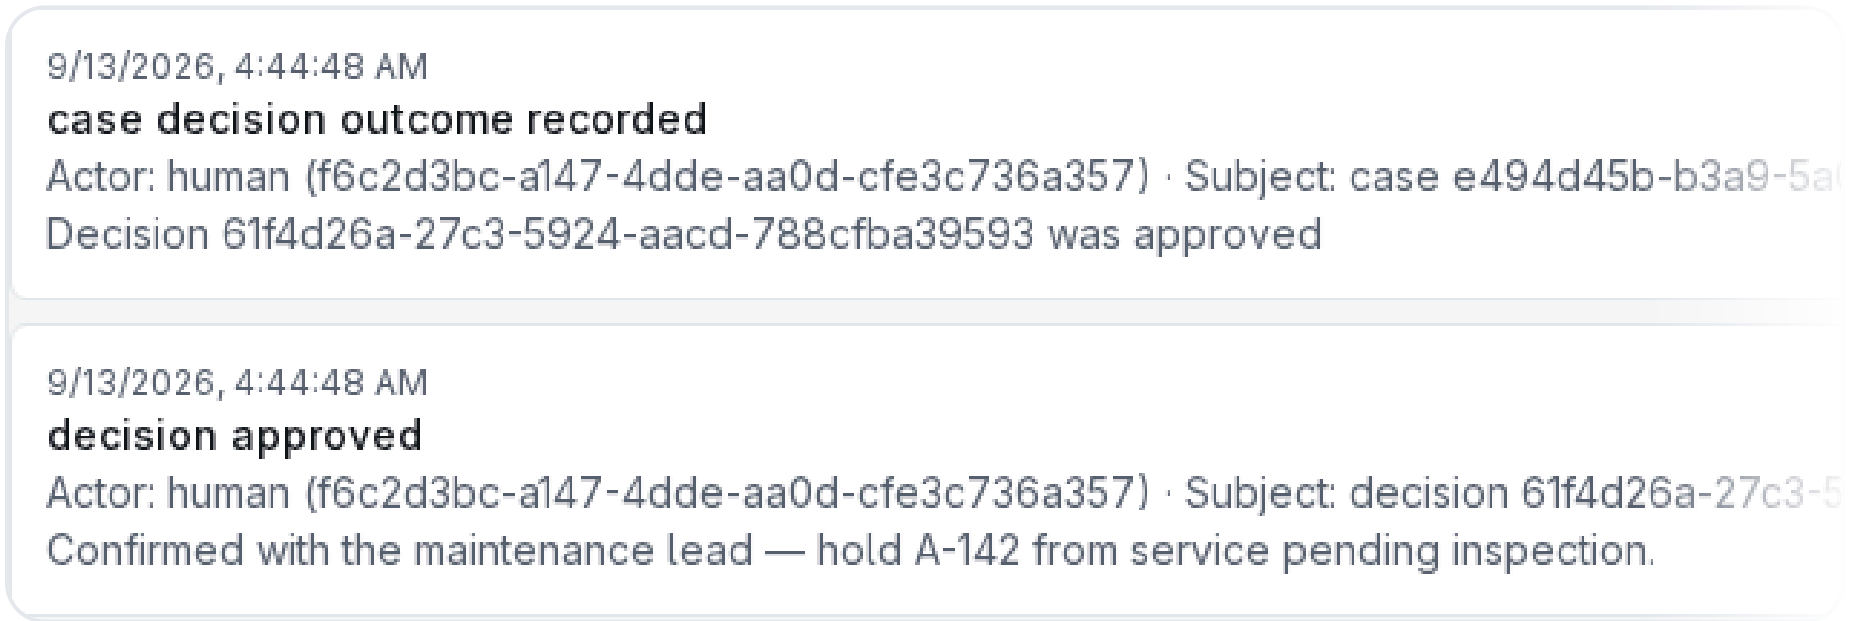
Offer more than accept and reject

Accept, Correct, Reject, Mark insufficient evidence, Link entity, Create entity, Add reviewer note. "The evidence is not enough" is a valid outcome.



Leave an audit entry

Accepting appends an audit entry. In the recorded case: pending to accepted, and one "observation accepted" entry.



Audit view: each entry shows actor, subject and timestamp.

What this does not show

Fixture data, a scripted review, no live model. The layout is the claim, not the accuracy. In this demo, the one vague 70% item is what reached the queue.